

## **Lesson Plan:**

### **Applications Development Laboratory (Machine Learning Applications)**

**Course Code** CS33002 (L-T-P-Cr: 0-0-4-2)

**Credits:** 2

**Session:** December-2024 to May-2025

**Course faculty:** Dr Partha Pratim Sarangi

**Email:** pp.sarangifcs@kiit.ac.in

#### **Week 1: Introduction to Machine Learning**

Objective: Understand the basics of machine learning and Python / R, etc

##### **Experiment:**

- Write a Python program to demonstrate data preprocessing steps: handling missing values, encoding categorical data, and feature scaling.
- Load a sample dataset using pandas (e.g., Iris or a custom dataset).
- Plot the distribution of a feature using `matplotlib.pyplot.hist()`.
- Create scatter plots to understand relationships between features using `seaborn.scatterplot()`.
- Use a correlation heatmap to find the relationship between multiple features with `seaborn.heatmap()`.

#### **Week 2: Regression Analysis**

Objective: Learn simple linear regression for predictive analysis.

##### **Experiment:**

Implement simple linear regression to predict house prices based on features such as area and number of rooms.

- Evaluate the model using metrics such as `mean_squared_error` and `r2_score`.

- Visualize the regression line on a scatter plot of the data

*Note: A common dataset we can explore is the "House Prices - Advanced Regression Techniques" dataset from kaggle. The UCI Machine Learning Repository has many datasets for classification, regression, and clustering tasks. A popular dataset related to housing is the "California Housing Prices" dataset.*

### **Week 3: Multiple Linear Regression**

Objective: Extend regression analysis to handle multiple predictors.

#### **Experiment:**

Implement multiple linear regression to predict student performance based on study hours, class attendance, and assignment scores.

- Calculate metrics like Mean Squared Error (MSE) and  $R^2$  Score to assess performance.
- Create a scatter plot comparing the actual vs. predicted values

*Note: We can get Student Performance Dataset from kaggle*

### **Week 4: Classification with Logistic Regression**

Objective: Understand binary classification using logistic regression.

#### **Experiment:**

Build a logistic regression model to classify emails as spam or not spam.

- Accuracy: A percentage indicating how many emails were correctly classified.
- Assess model performance using accuracy, confusion matrix, and classification report (Classification Report: Precision, Recall, and F1-score for both spam and non-spam classes)

*Note: We will use a publicly available dataset such as the SpamBase dataset from the UCI Machine Learning Repository or create a simple synthetic dataset for demonstration purposes*

## **Week 5: Decision Trees**

Objective: Learn decision tree classifiers for non-linear decision boundaries.

### **Experiment:**

Develop a decision tree model to classify species in the *Iris dataset*.

- Evaluate the model using accuracy and a detailed classification report with precision, recall, and F1-score for each species.
- Use `plot_tree()` to visualize the decision tree. This helps understand how the model splits the data at different nodes.

## **Week 6: Random Forests**

Objective: Understand ensemble methods using random forests.

### **Experiment:**

Implement a random forest classifier to predict customer churn in a telecom dataset.

- Accuracy: Find overall prediction accuracy.
- Confusion Matrix: Breakdown of true/false positives and negatives.
- Classification Report: Includes precision, recall, and F1-score for both churned and retained classes.

*Note: You can use a real-world dataset such as the Telco Customer Churn Dataset available on Kaggle or create a similar dataset for practice.*

## **Week 7: Support Vector Machines (SVM)**

Objective: Explore SVM for linear and non-linear classification.

### **Experiment:**

Build an SVM model to classify handwritten digits using the MNIST dataset.

- Accuracy: Overall percentage of correct predictions.
- Confusion Matrix: Provides a breakdown of correct/incorrect predictions for each digit.
- Classification Report: Detailed precision, recall, and F1-score for each digit.

*Note: The MNIST dataset consists of 70,000 grayscale images of handwritten digits (28x28 pixels) representing numbers from 0 to 9. We can use `sklearn.datasets.fetch_openml` to load the dataset.*

## **Week 8: K-Means Clustering**

Objective: Learn unsupervised learning through clustering.

### **Experiment:**

Apply K-Means clustering to group customers based on purchasing patterns.

- Plot the clusters in 2D space (using scaled features) and mark the centroids.
- Analyze each cluster to identify customer groups based on purchasing patterns, such as:
  - High-income, high-spending customers.
  - Low-income, low-spending customers

*Note: We can use the Mall Customers Dataset. This dataset is available in .csv format, typically containing columns such as CustomerID, Gender, Age, Annual Income, and Spending Score (1-100).*

## **Week 9: Principal Component Analysis (PCA)**

Objective: Reduce dimensionality using PCA.

### **Experiment:**

Perform PCA on a high-dimensional dataset and visualize the results.

- Displays the proportion of variance captured by each principal component.
- 2D Plot: Projects data onto the first two principal components.
- 3D Plot: Projects data onto the first three principal components.

*Note: We can use the Wine Dataset from the sklearn library. It has 13 features (high dimensionality) related to chemical properties of wine samples, which can be reduced to 2 or 3 principal components for visualization.*

## **Week 10: Neural Networks with TensorFlow/Keras**

Objective: Introduce deep learning basics with neural networks.

### **Experiment:**

Build a simple neural network to classify handwritten digits using TensorFlow/Keras.

- The images are flattened from 28x28 matrices into 1D vectors of length 784 to be fed into the neural network.
- Input Layer: The input is a flattened vector of size 784 (28x28).
- Hidden Layer: A dense layer with 128 neurons and ReLU activation. This layer allows the model to learn complex patterns.
- Dropout Layer: A dropout layer with a rate of 0.2, which helps reduce overfitting by randomly setting some of the weights to zero during training.

- Output Layer: A softmax layer with 10 neurons for classifying the digits from 0 to 9.
- Show the Test Accuracy, model's performance over epochs (accuracy plot), loss plot indicates whether the model is converging

*Note: We can use MNIST Dataset: The MNIST dataset contains 60,000 training images and 10,000 test images of handwritten digits (0-9).*

## **Week 11: Convolutional Neural Networks (CNN)**

Objective: Learn image classification using CNNs.

### **Experiment:**

Implement a CNN model to classify images from the CIFAR-10 dataset.

- Show the Test Accuracy, Training vs Validation Accuracy, Training vs Validation Loss and Predictions

*Note: CIFAR-10 Dataset: The CIFAR-10 dataset contains 60,000 32x32 color images in 10 classes. CIFAR-10 dataset can be loaded using `cifar10.load_data()`.*

**Capstone Application** ( Any of the application must be deployed on the web and accessible through a browser)

Objective: Integrate learned concepts into a real-world application.

### **Experiment:**

Hosting an ML application on the web with a front end that can be run in a browser including creating the machine learning model, developing a web interface (front end), setting up a back-end server, and hosting the entire application on the web.

-----

**Create a Web Front-End (User Interface):**

Languages: Use HTML, CSS, and JavaScript (or frameworks like React, Angular, or Vue.js) to create the user interface.

### **Develop the Back-End (Server):**

Flask or FastAPI (for Python): These are lightweight web frameworks for serving machine learning models.

### **Connect the Front-End to the Back-End:**

Use AJAX or Fetch API in JavaScript to send user input to the back-end (Flask/FastAPI)

### **Learning Outcomes:**

- Understand and apply various supervised and unsupervised machine learning algorithms.
- Build and evaluate models using Python and its libraries (e.g., NumPy, pandas, scikit-learn, TensorFlow).
- Solve real-world problems using machine learning techniques.

### **Tools Required:**

- Python 3.x/ R
- Libraries: NumPy, pandas, matplotlib, scikit-learn, TensorFlow/Keras, NLTK/Spacy